

Encrypted Traffic Classification: From Small Data to Scalable Production

Methods for Detecting Cyber Attacks

Raz Cohen, Dael Hacohen Waingarten, Shir Bismuth

The ‘Small Data’ Challenge: From Big Data to Smart Data

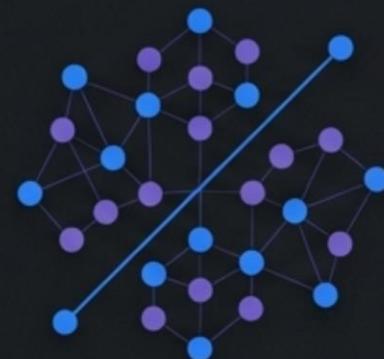
Identifying 128 applications from sparse samples required a strategic shift from ‘Big Data’ volume to ‘Smart Data’ intelligence.

The Challenge: High Sparsity, High Cardinality



- **128 Applications:** Massive class imbalance due to a ‘long tail’ effect with few training examples per class.
- **5 Attribution Types:** Limited instances to differentiate nuanced behaviors like real-time audio vs. video.
- **Encrypted Payloads:** Traditional port-based or deep-packet inspection methods are rendered obsolete.

Our Strategy: Engineering Resilience



- **Class-Weight Balancing:** Configured models with `class_weight='balanced'` to force equal attention to rare classes.
- **Domain-Driven Feature Selection:** Engineered features based on network protocol theory to capture stable signals and reduce noise.

Feature Engineering: Uncovering the Behavioral DNA of Encrypted Flows



Application Identification (128 Classes)

Key Insight: Focused on the unencrypted metadata of the TLS negotiation phase.

Top Features

`handshake_avg` & `handshake_std`
Measures timing and size patterns of the initial TLS handshake. The negotiation protocol itself is a unique signature.

`fwd_bwd_pkts_diff`
Captures flow asymmetry, distinguishing upload-heavy vs. download-heavy applications.



Attribution (5 Classes)

Key Insight: Created custom signatures to model spatio-temporal usage patterns.

Custom Signatures

`chat_signature`
A metric combining `silence_ratio` and `burstiness` to identify messaging patterns (long silences, short data bursts).

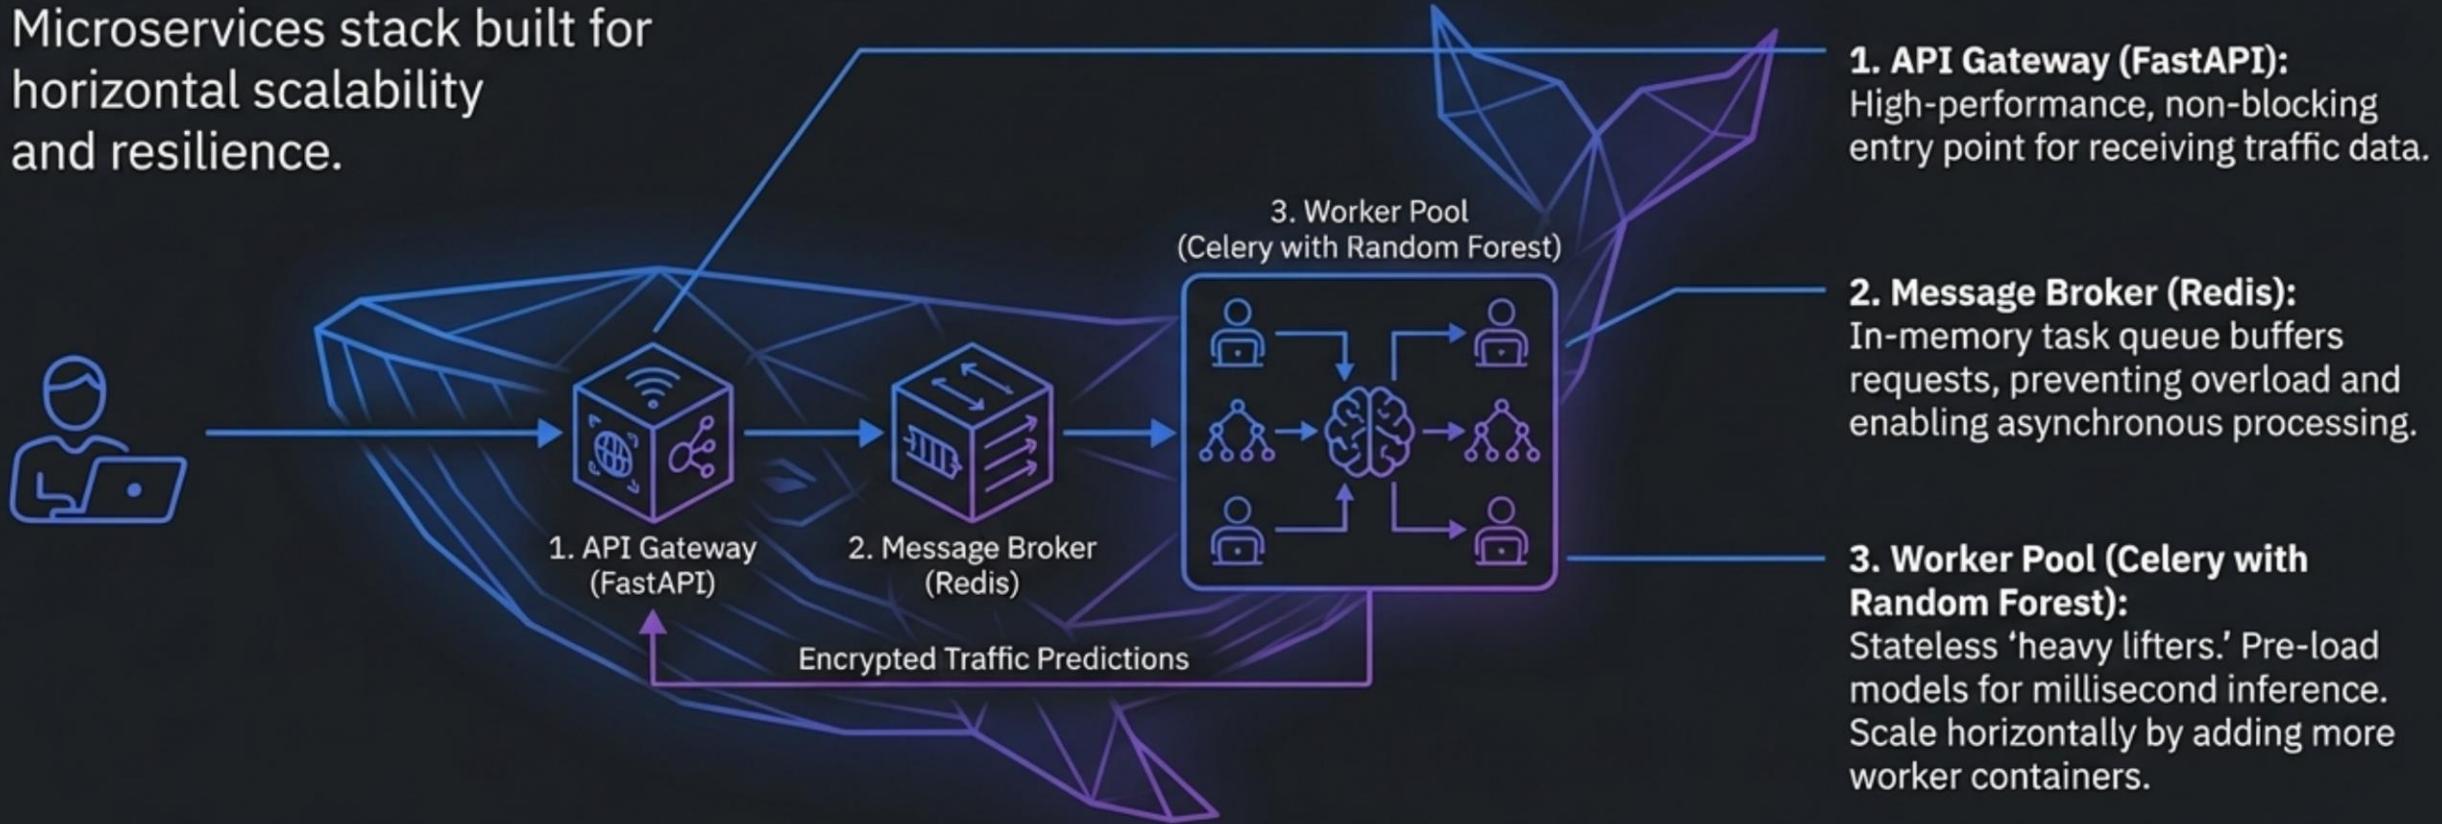


`stream_signature`
Highlights high-volume, high-throughput, low-silence traffic typical of VOD or large file downloads.



System Architecture: Engineered for the Hexabyte Era

A decoupled, asynchronous Microservices stack built for horizontal scalability and resilience.



This design ensures the UI remains responsive and inference can scale to handle massive telecom traffic without changing core code.

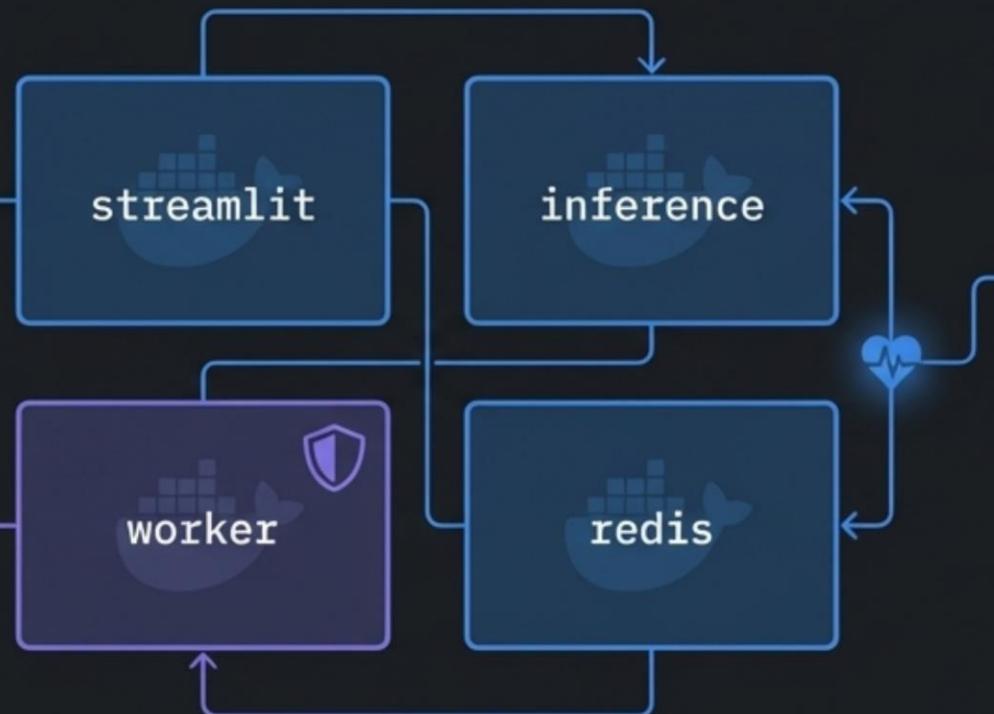
Secure & Reliable Deployment: A Containerized Approach

Multi-Container Orchestration

- We use Docker Compose to manage four specialized services: `streamlit`, `inference`, `worker`, and `redis`.
- This separation optimizes image sizes and allows independent scaling of the ML workload.

Embedded Cybersecurity Practices

- **Non-Root Execution:** Containers run under a dedicated `appuser`, preventing root access to the host.
- **Minimal Attack Surface:** We use `python:3.11-slim` base images to reduce vulnerabilities and storage overhead.



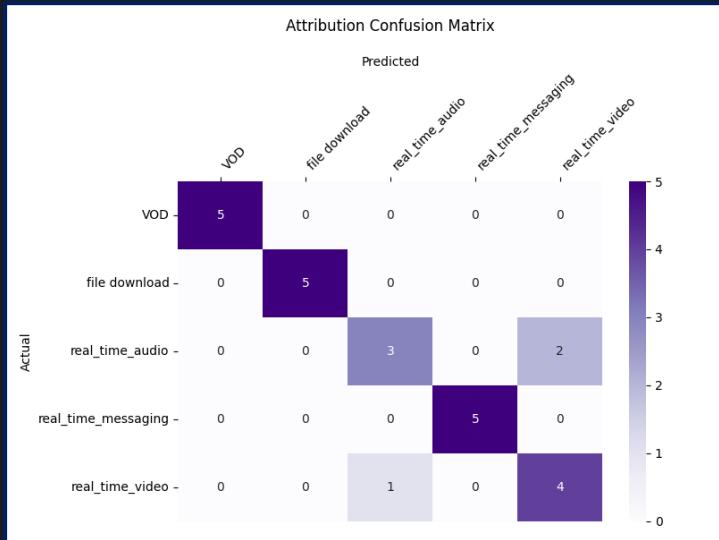
Production-Grade Reliability

- **Health Checks:** System ensures the Redis broker is 'Healthy' before starting dependent services, preventing startup failures.
- **Environment Consistency:** Docker guarantees an identical environment from development to production, eliminating 'it works on my machine'.

Results & Impact: Depth, Accuracy, and Scale

Attribution Classification (5 Classes)

88% Accuracy



- Perfect classification for VOD, File Download, and Messaging due to distinct behavioral signatures.
- Primary challenge: confusion between `real_time_audio` and `real_time_video` due to similar low-latency UDP patterns.

Application Identification (128 Classes)

51.7% Accuracy
66x Better Than a Random Guess

(0.78% baseline)

- Model excels at identifying apps with unique TLS ‘handshake’ dynamics, proving our ‘Behavioral DNA’ approach.
- Example apps with high F1-scores:
“brainfund” (1.00), ‘vimeo’ (1.00), ‘ndtv’ (1.00), ‘wprdc.org’ (1.00)